
Kubernetes Deployment + Service (Best Practice 2025)

1. QUICK ONE-LINERS (perfect for labs, testing, demos)

Create a Deployment with 3 replicas

```
kubectl create deployment nginx-deploy --image=nginx:alpine --replicas=3
```

Expose it internally (ClusterIP)

```
kubectl expose deployment nginx-deploy --port=80 --name=nginx-svc
```

Expose it externally (NodePort)

```
kubectl expose deployment nginx-deploy --port=80 --type=NodePort --name=nginx-svc
```

```
# Expose with LoadBalancer (GKE, EKS,  
AKS, DigitalOcean, etc.)
```

```
kubectl expose deployment nginx-deploy -  
-port=80 --type=LoadBalancer --  
name=nginx-lb
```

```
# Check everything
```

```
kubectl get deploy,rs,pods,svc
```

```
kubectl get svc nginx-svc # note the  
NodePort or External IP
```

```
# Access (NodePort example)
```

```
# http://<any-node-ip>:<node-port> e.g.,  
http://192.168.1.100:31567
```

```
# Scale instantly
```

```
kubectl scale deployment nginx-deploy --  
replicas=10
```

```
kubectl scale deployment nginx-deploy --  
replicas=2
```

```
#  
=====  
=====  
# BEST PRACTICE: Proper YAML files  
(production-ready)  
#  
=====  
=====
```



```
# File: deployment.yaml  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: nginx-deploy  
  labels:  
    app: nginx  
spec:  
  replicas: 3
```

selector:

matchLabels:

app: nginx

template:

metadata:

labels:

app: nginx

spec:

containers:

 - **name: nginx**

image: nginx:alpine

ports:

 - **containerPort: 80**

resources:

requests:

memory: "64Mi"

cpu: "100m"

limits:

```
memory: "128Mi"
cpu: "200m"

# File: service.yaml
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: NodePort          # Use ClusterIP,
NodePort, or LoadBalancer
```

```
# Apply both
```

```
kubectl apply -f deployment.yaml
```

```
kubectl apply -f service.yaml
```

```
# Upgrade image (zero downtime)
```

```
kubectl set image deployment/nginx-deploy nginx=nginx:1.25 --record
```

```
kubectl rollout status deployment/nginx-deploy
```

```
# Rollback if something goes wrong
```

```
kubectl rollout undo deployment/nginx-deploy
```

=====

ONE-LINE SUPER FAST VERSION (most people use this)

=====

```
kubectl create deployment myapp --  
image=nginx:alpine --replicas=4  
  
kubectl expose deployment myapp --  
type=NodePort --port=80 --name=myapp-  
svc
```

```
# Done! Access via the NodePort shown  
in:
```

```
kubectl get svc myapp-svc
```

```
=====
```

```
CLEANUP
```

```
=====
```

```
kubectl delete deployment nginx-deploy  
kubectl delete service nginx-service
```

```
# or
```

```
kubectl delete -f deployment.yaml -f  
service.yaml
```

=====

SUMMARY: This is the correct way in
2025

=====

Deployment = modern, self-healing, rolling
updates, rollback

ReplicaSet = created automatically by
Deployment

ReplicationController = old/legacy, don't
use

Pod alone = no scaling or healing

Use Deployment + Service → always!

Done! You now have a production-grade,
scalable, updatable Nginx app running in
Kubernetes.